Nikos Litharis

AI Engineer — LLM Systems, RAG & Evaluation

nicklitharis@gmail.com | +30 697 658 4914 | Thessaloniki, Greece | github.com/NickLitharis | linkedin.com/in/nicklitharis

SUMMARY

AI Engineer building production LLM systems in Python on Azure — RAG, multi-agent orchestration and evaluation pipelines — for public-sector clients including Greece's national power transmission operator (IPTO) and the Independent Authority for Public Revenue (AADE), building TAXIS, a legal-research assistant on the Microsoft Agent Framework.

WORK EXPERIENCE

Netcompany **Nov 2025 – Present**

AI Engineer

- Owned the evaluation and UAT pipeline for TAXIS, AADE's production legal-research assistant, end to end: versioned scenario datasets covering adversarial, citation, temporal and incomplete-context challenge sets, with gold sources, a CLI runner that scores answers and citations, and evidence-promotion tooling with runtime provenance.
- Designed the evaluation suite behind retrieval improvements in IPTO's RAG knowledge-base assistant: 50 bilingual (Greek/English) cases over IEC standards documentation, used to iterate on prompts, chunking and embeddings.
- Architected a durable, resumable batch-ingestion pipeline for large legal corpora (OpenAI Batch API for embeddings and enrichment, Azure AI Search v3), with recovery for expired leases and concurrent admission, and a live dashboard for throughput and queue flow.
- Added voice interaction to IPTO's RAG assistant by integrating Azure Speech SDK STT/TTS with managed identity authentication, handling long-response streaming, timeouts and cancellation.
- Built Docker-containerized multi-agent (A2A) orchestration pipelines with LangChain, reused across delivery projects to cut manual coordination overhead.

Military Service, Hellenic Army **Sep 2024 – Jun 2025**

The Prosperty **Feb 2024 – Mar 2024**

AI Development Intern

- Cut listing-description production time from ~2 hours to under 30 seconds by building a GPT-3 generation tool in Python, adopted for new property listings.
- Raised generated-copy acceptance by fine-tuning language models on property data and iterating on training samples against tone and accuracy requirements.
- Integrated the tool into the live listing platform by shipping a REST API with the product and engineering teams.

PROJECTS

Haystack & Elasticsearch Q&A Search — github.com/NickLitharis/haystack-elastic
University project with a collaborator: indexed scraped Q&A data into Elasticsearch and compared two retrieval pipelines — keyword (BM25) against semantic search with Haystack's EmbeddingRetriever, returning the top-5 answers by embedding similarity.

Earthquake Analysis and Prediction — github.com/NickLitharis/earthquake-prediction
Predicted earthquake magnitude from 251,000 historical seismic records by building an end-to-end ML pipeline (preprocessing, feature engineering, Scikit-Learn), comparing three regression models and reaching R^2 0.86 with Gradient Boosting.

MedLeaf RAG *(in progress)*
Bilingual (Greek/English) retrieval-augmented QA over official EU medicine information, with cited answers, medical safety guardrails, and a CI evaluation gate that blocks regressions.

EDUCATION

BSc in Computer Science, Ionian University, Corfu **2019 – 2024**
Thesis: Text Generation from Medical Images — CNN encoder with an attention-based RNN decoder in TensorFlow, generating captions for medical images.

SKILLS

GenAI / LLM: RAG, Prompt Engineering, Embeddings, LLM Evaluation, Fine-Tuning, Multi-Agent Orchestration (A2A), Microsoft Agent Framework, LangChain, Azure OpenAI, GPT APIs, Azure AI Search, PostgreSQL Vector Search

Programming: Python (NumPy, Pandas, Scikit-Learn, Keras), C# / .NET, TypeScript / Angular, SQL (Oracle, PostgreSQL)

ML Engineering: Docker, Docker Compose, GitLab CI, Git, Azure, REST APIs, Automated Regression Testing (.NET, Angular, Python), LLM Evaluation Harnesses, Model Monitoring

AI-Assisted Dev: Claude Code, Pi (multi-model: GPT, DeepSeek), agentic coding workflows